

# How to build a reliable, highly available OpenStack system?

Attila Szlovencsák

CTO, Component Soft  
[attila@componentsoft.eu](mailto:attila@componentsoft.eu)




# Pets or cattle?



- Pet

- has loving names
- unique, hand raised
- cared for, nursed when get ill
- it gets bigger over time



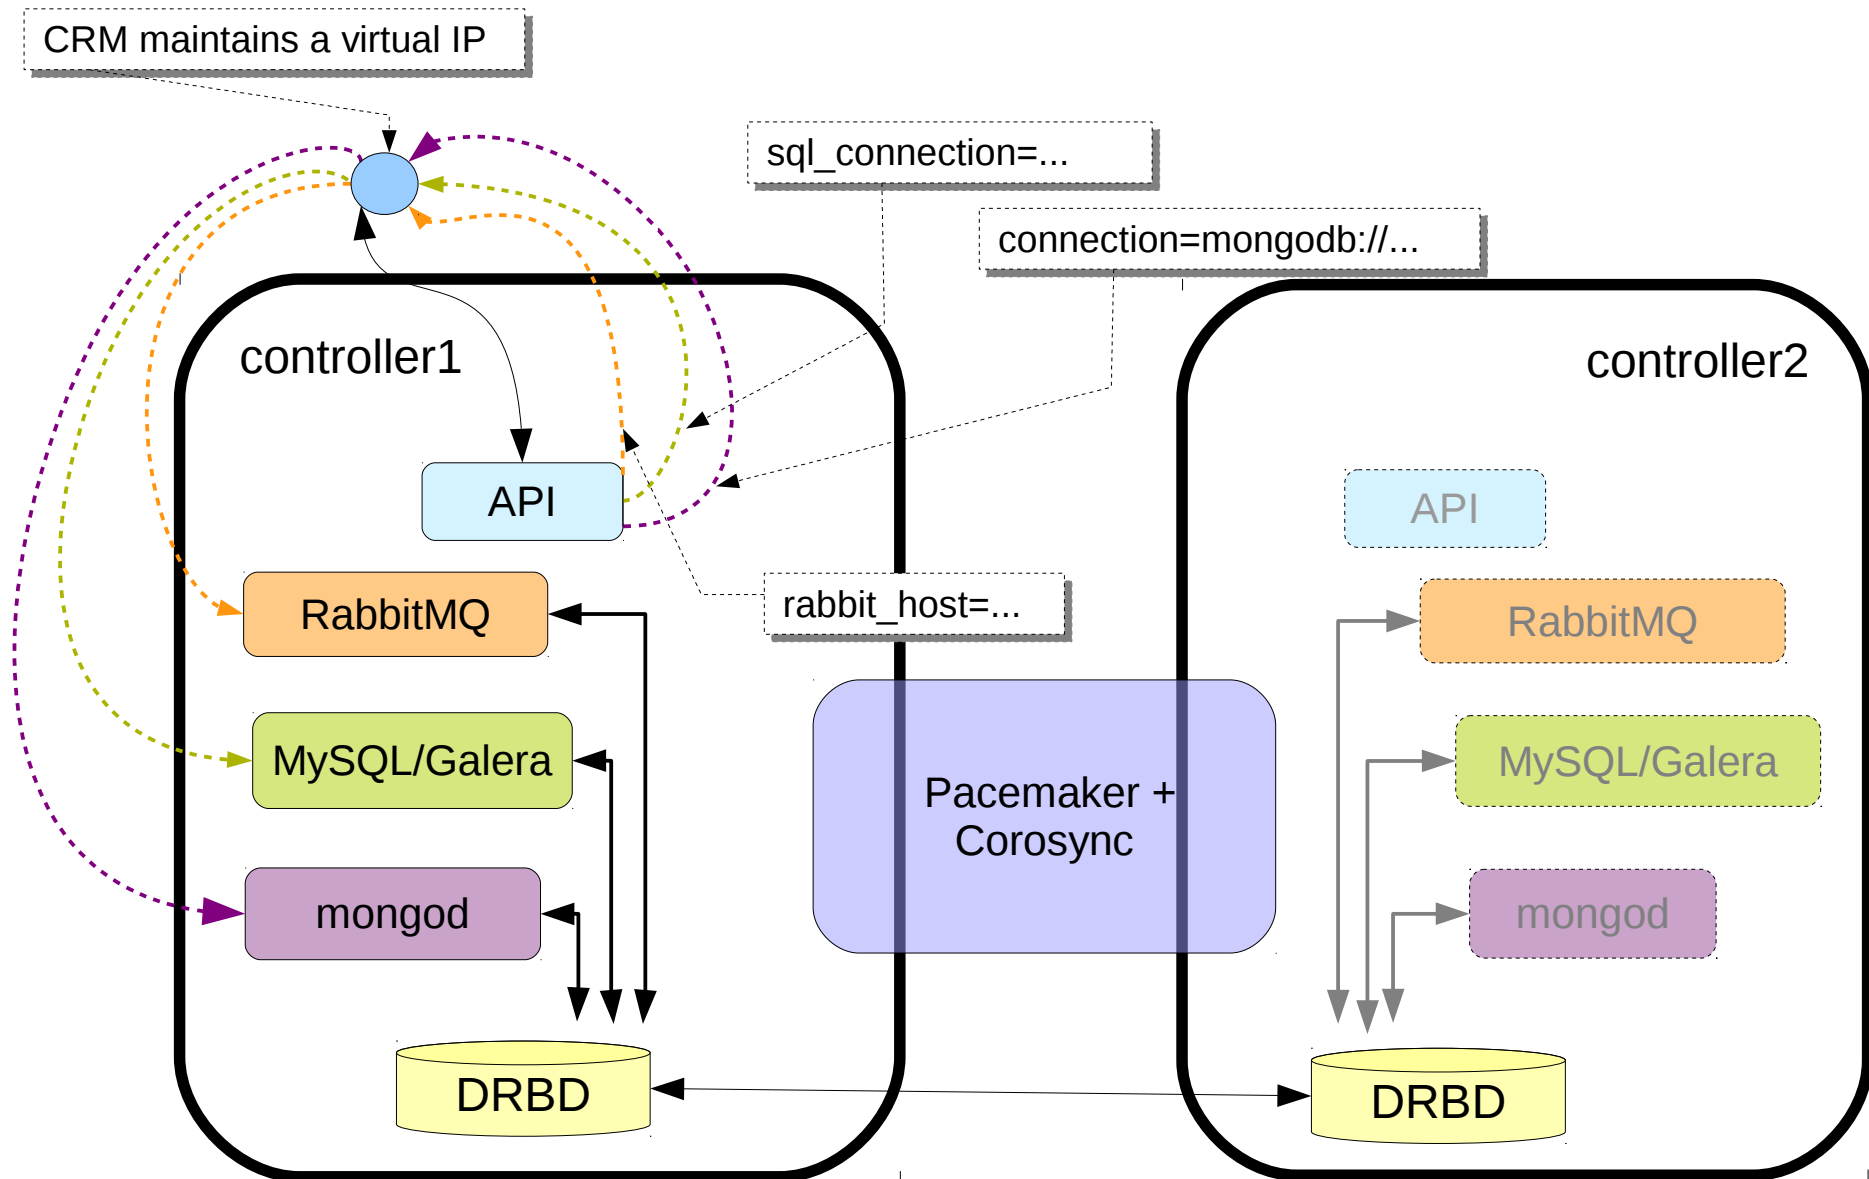
- Cattle

- has a number (web001)
- almost identical to another one
- when they get sick, you shoot them, and buy another one
- you get more over time

We need *cattle* style VMs in the cloud



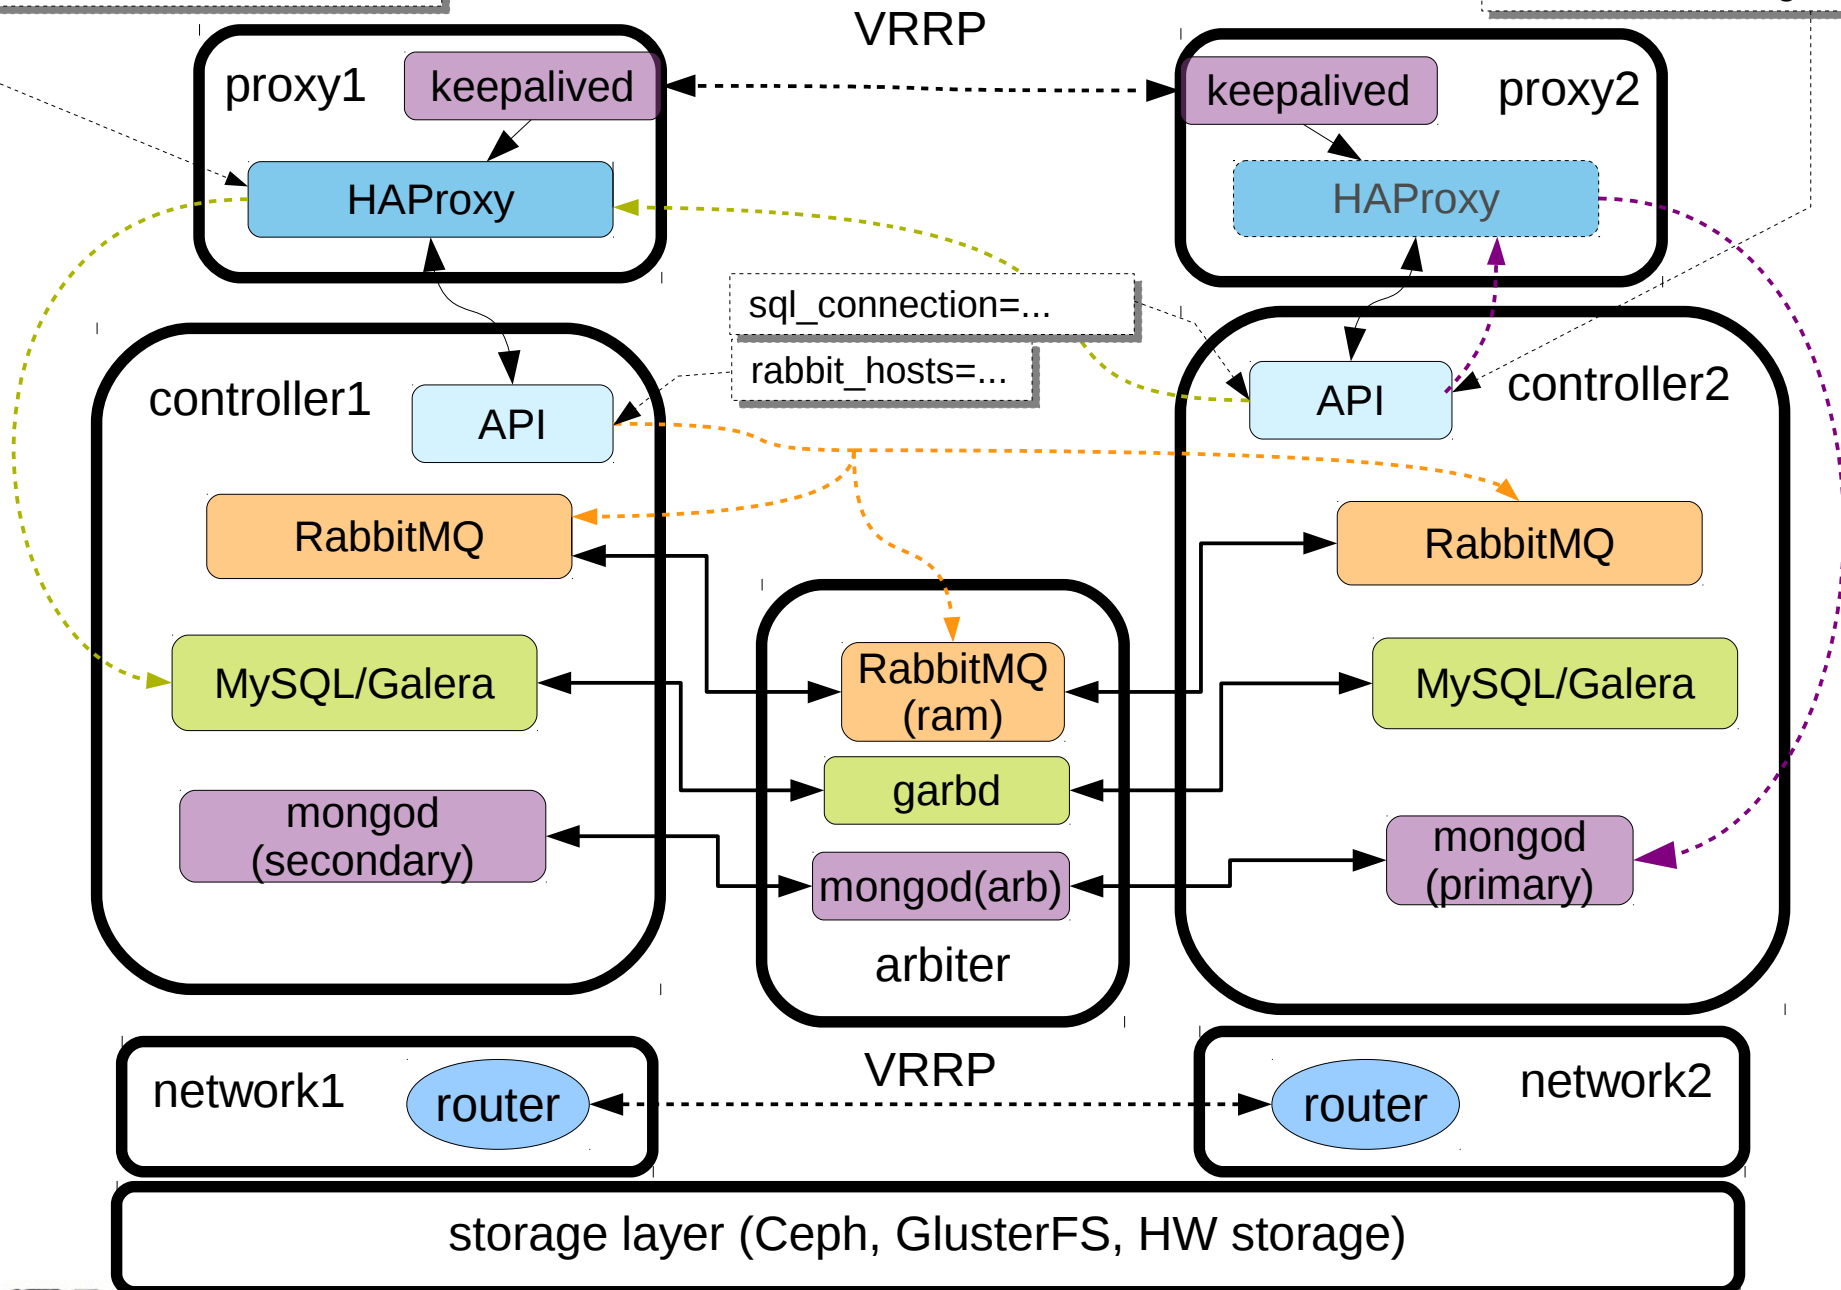
# Active/passive HA



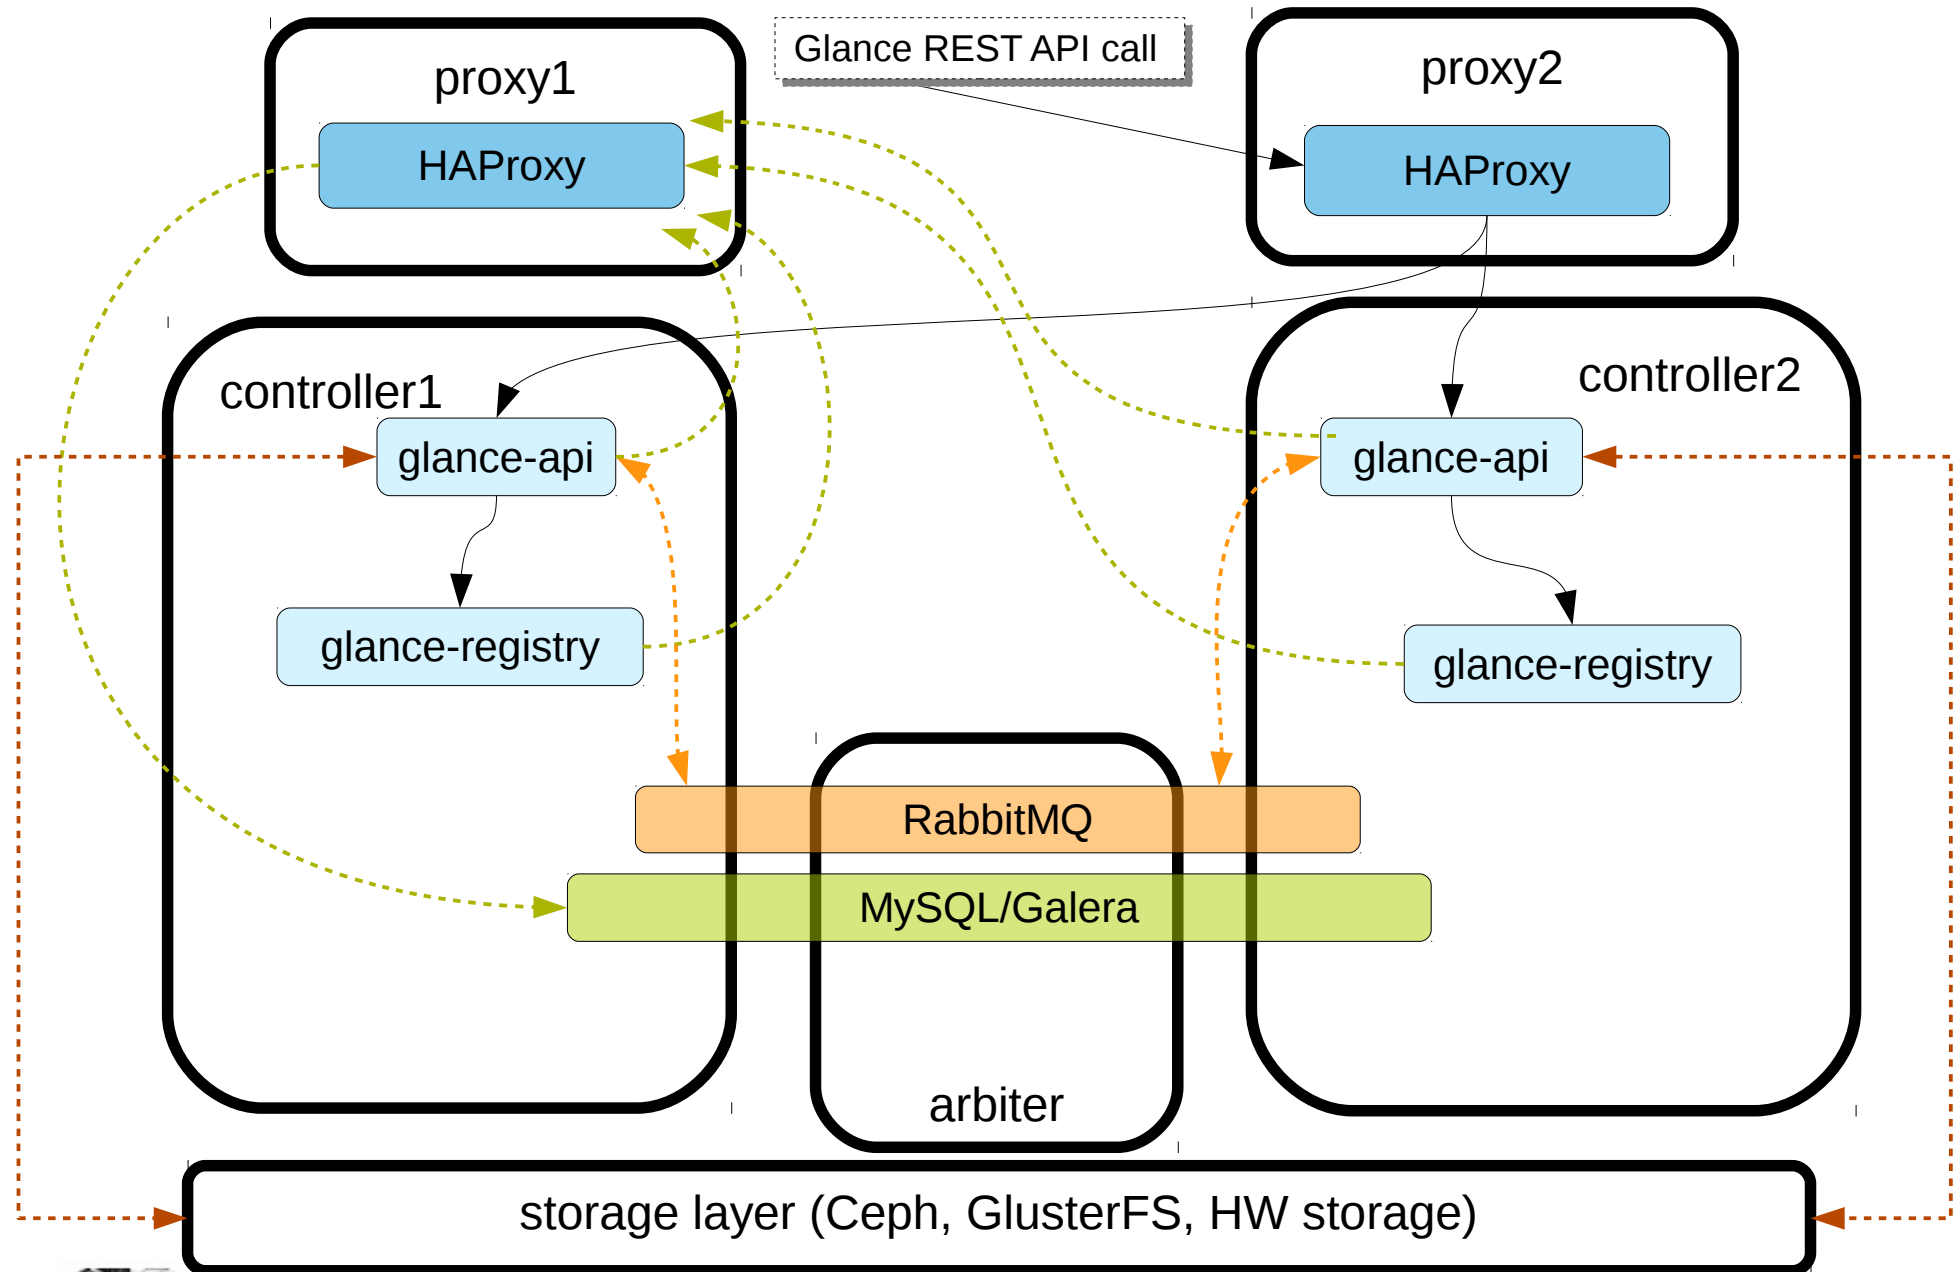
Active/backup config  
for Galera  
(at least for Neutron)

# Active/active HA

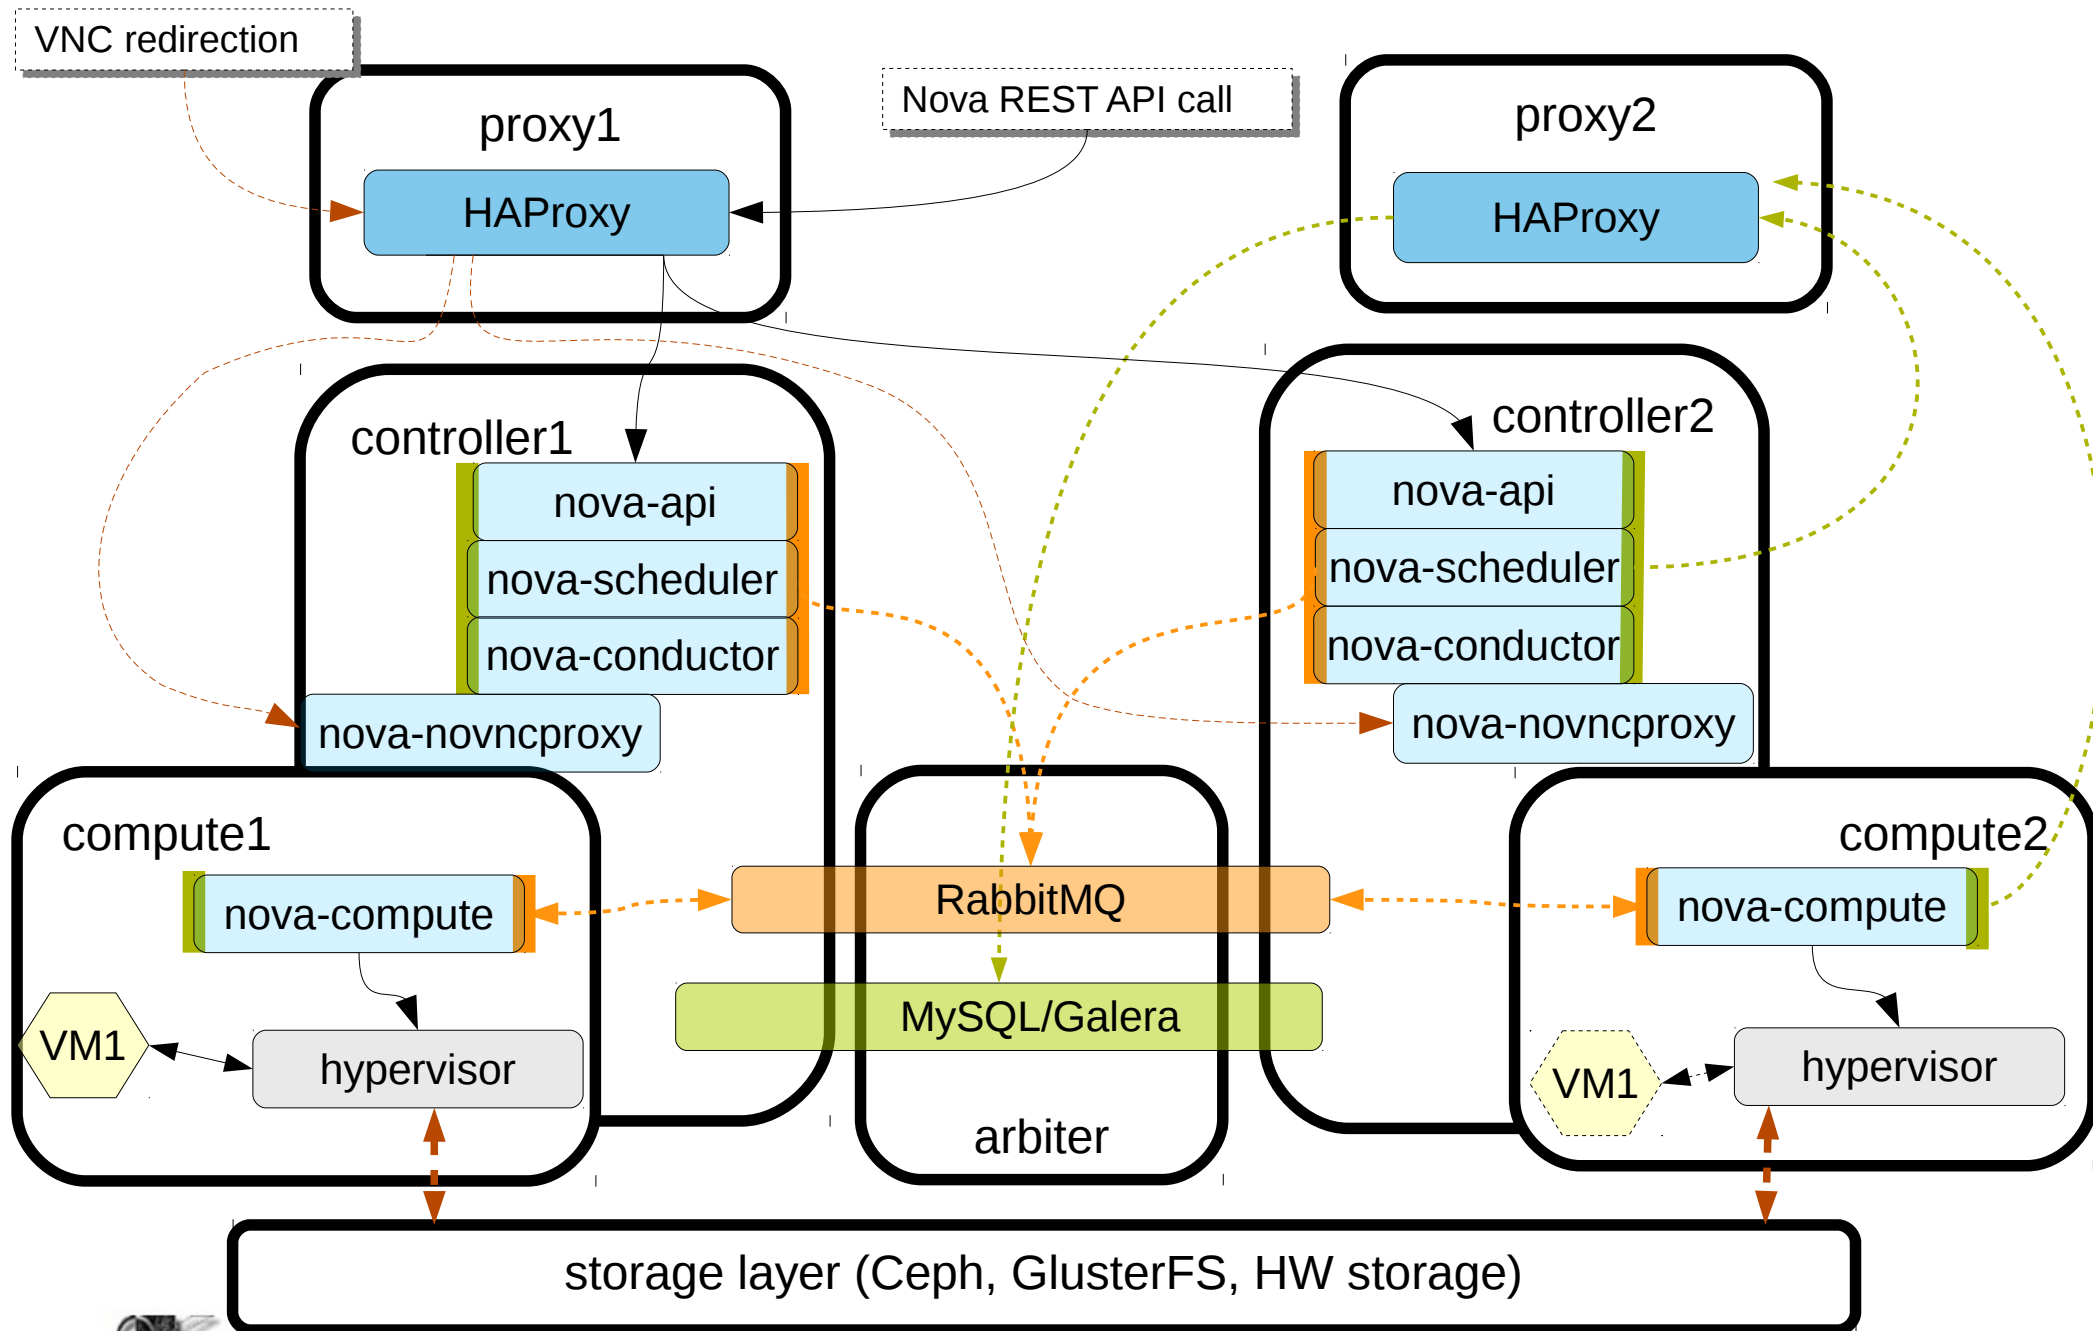
connection=mongodb://...



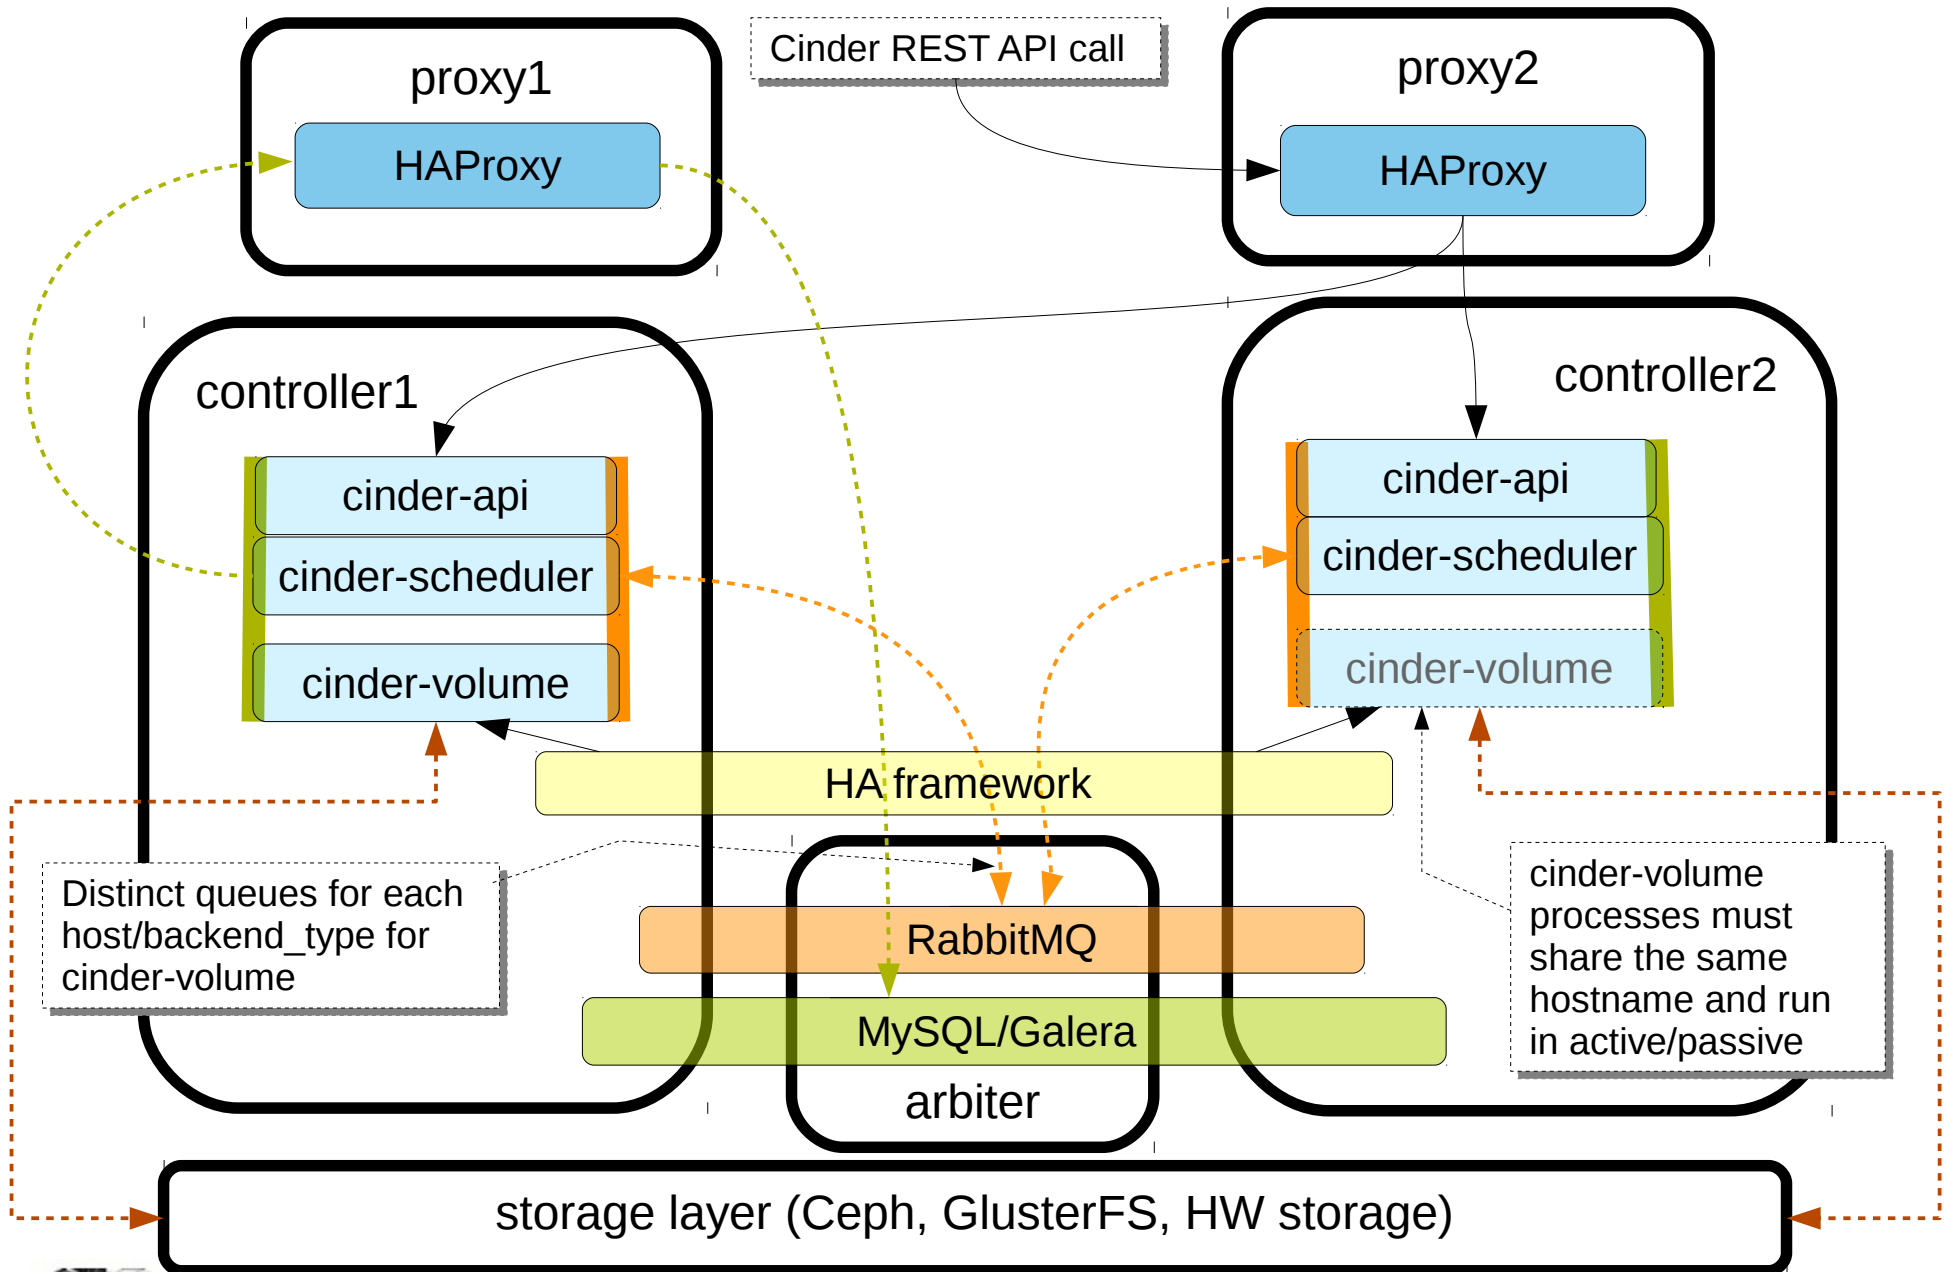
# Glance HA (active-active)



# Nova HA (active-active)

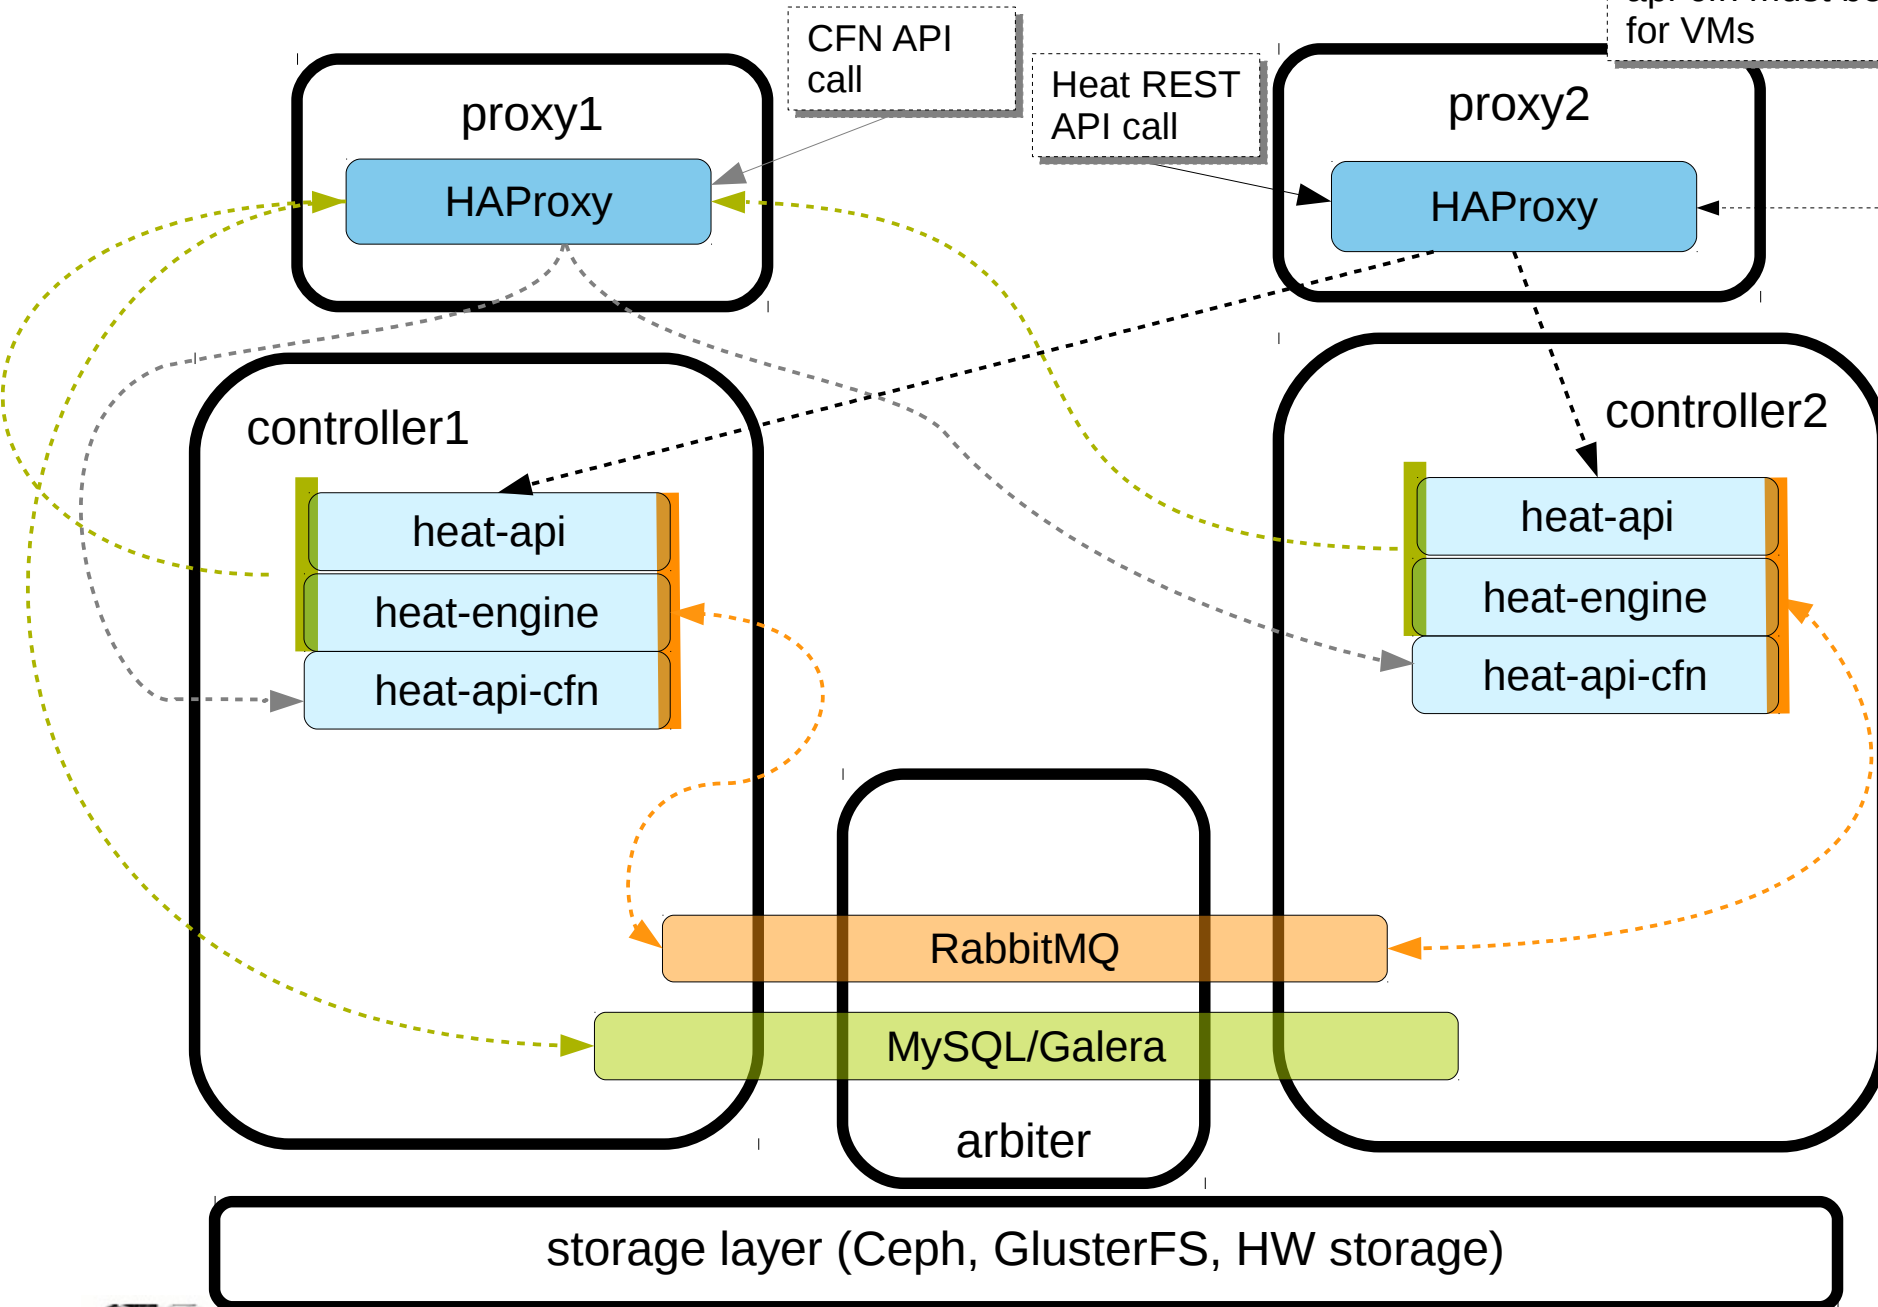


# Cinder HA (active-passive)\*



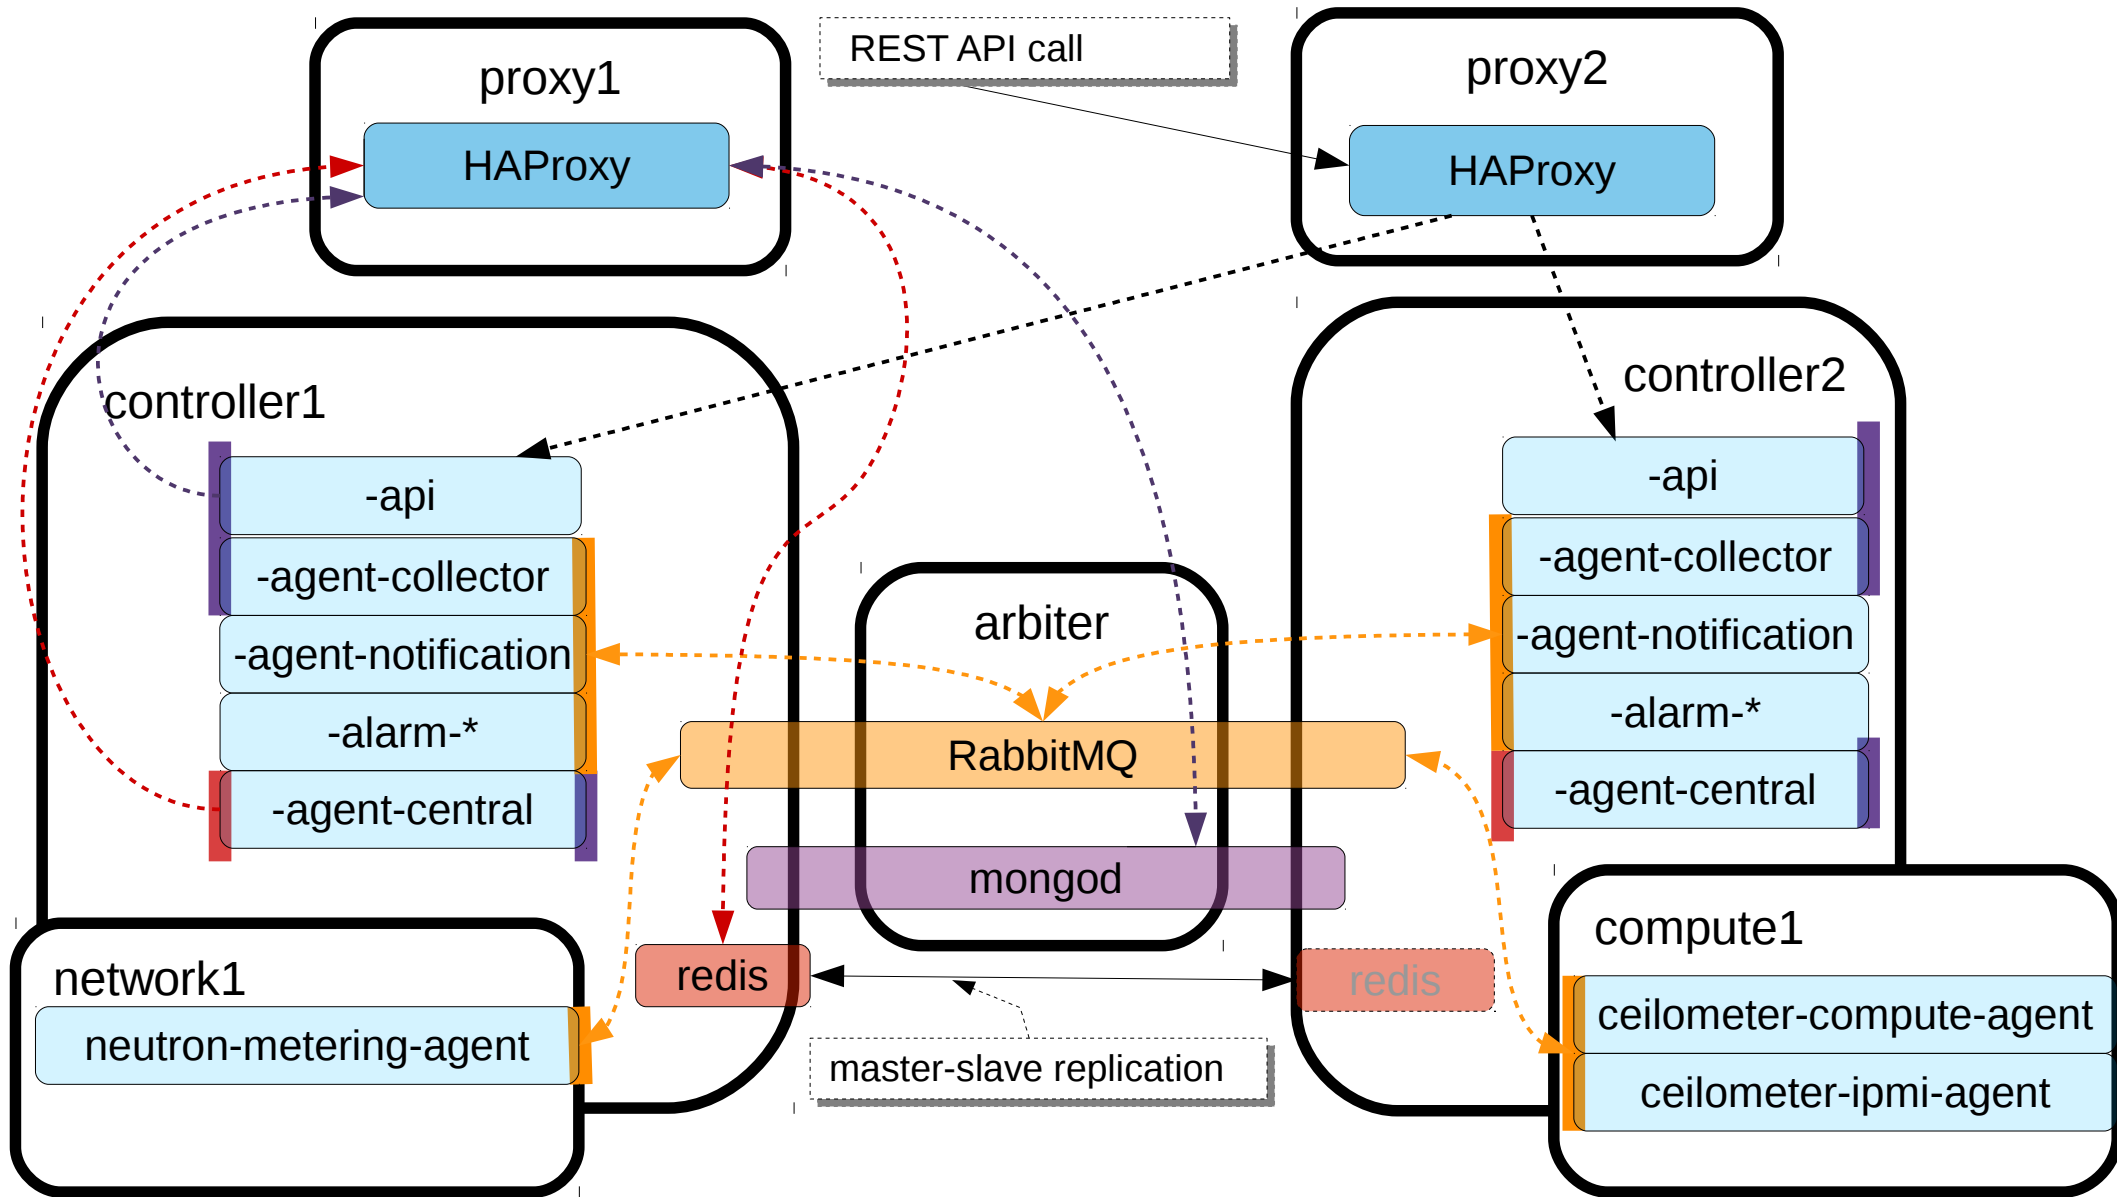
# Heat HA (active-active)

Both heat-api and heat-api-cfn needs LB. heat-api-cfn must be available for VMs

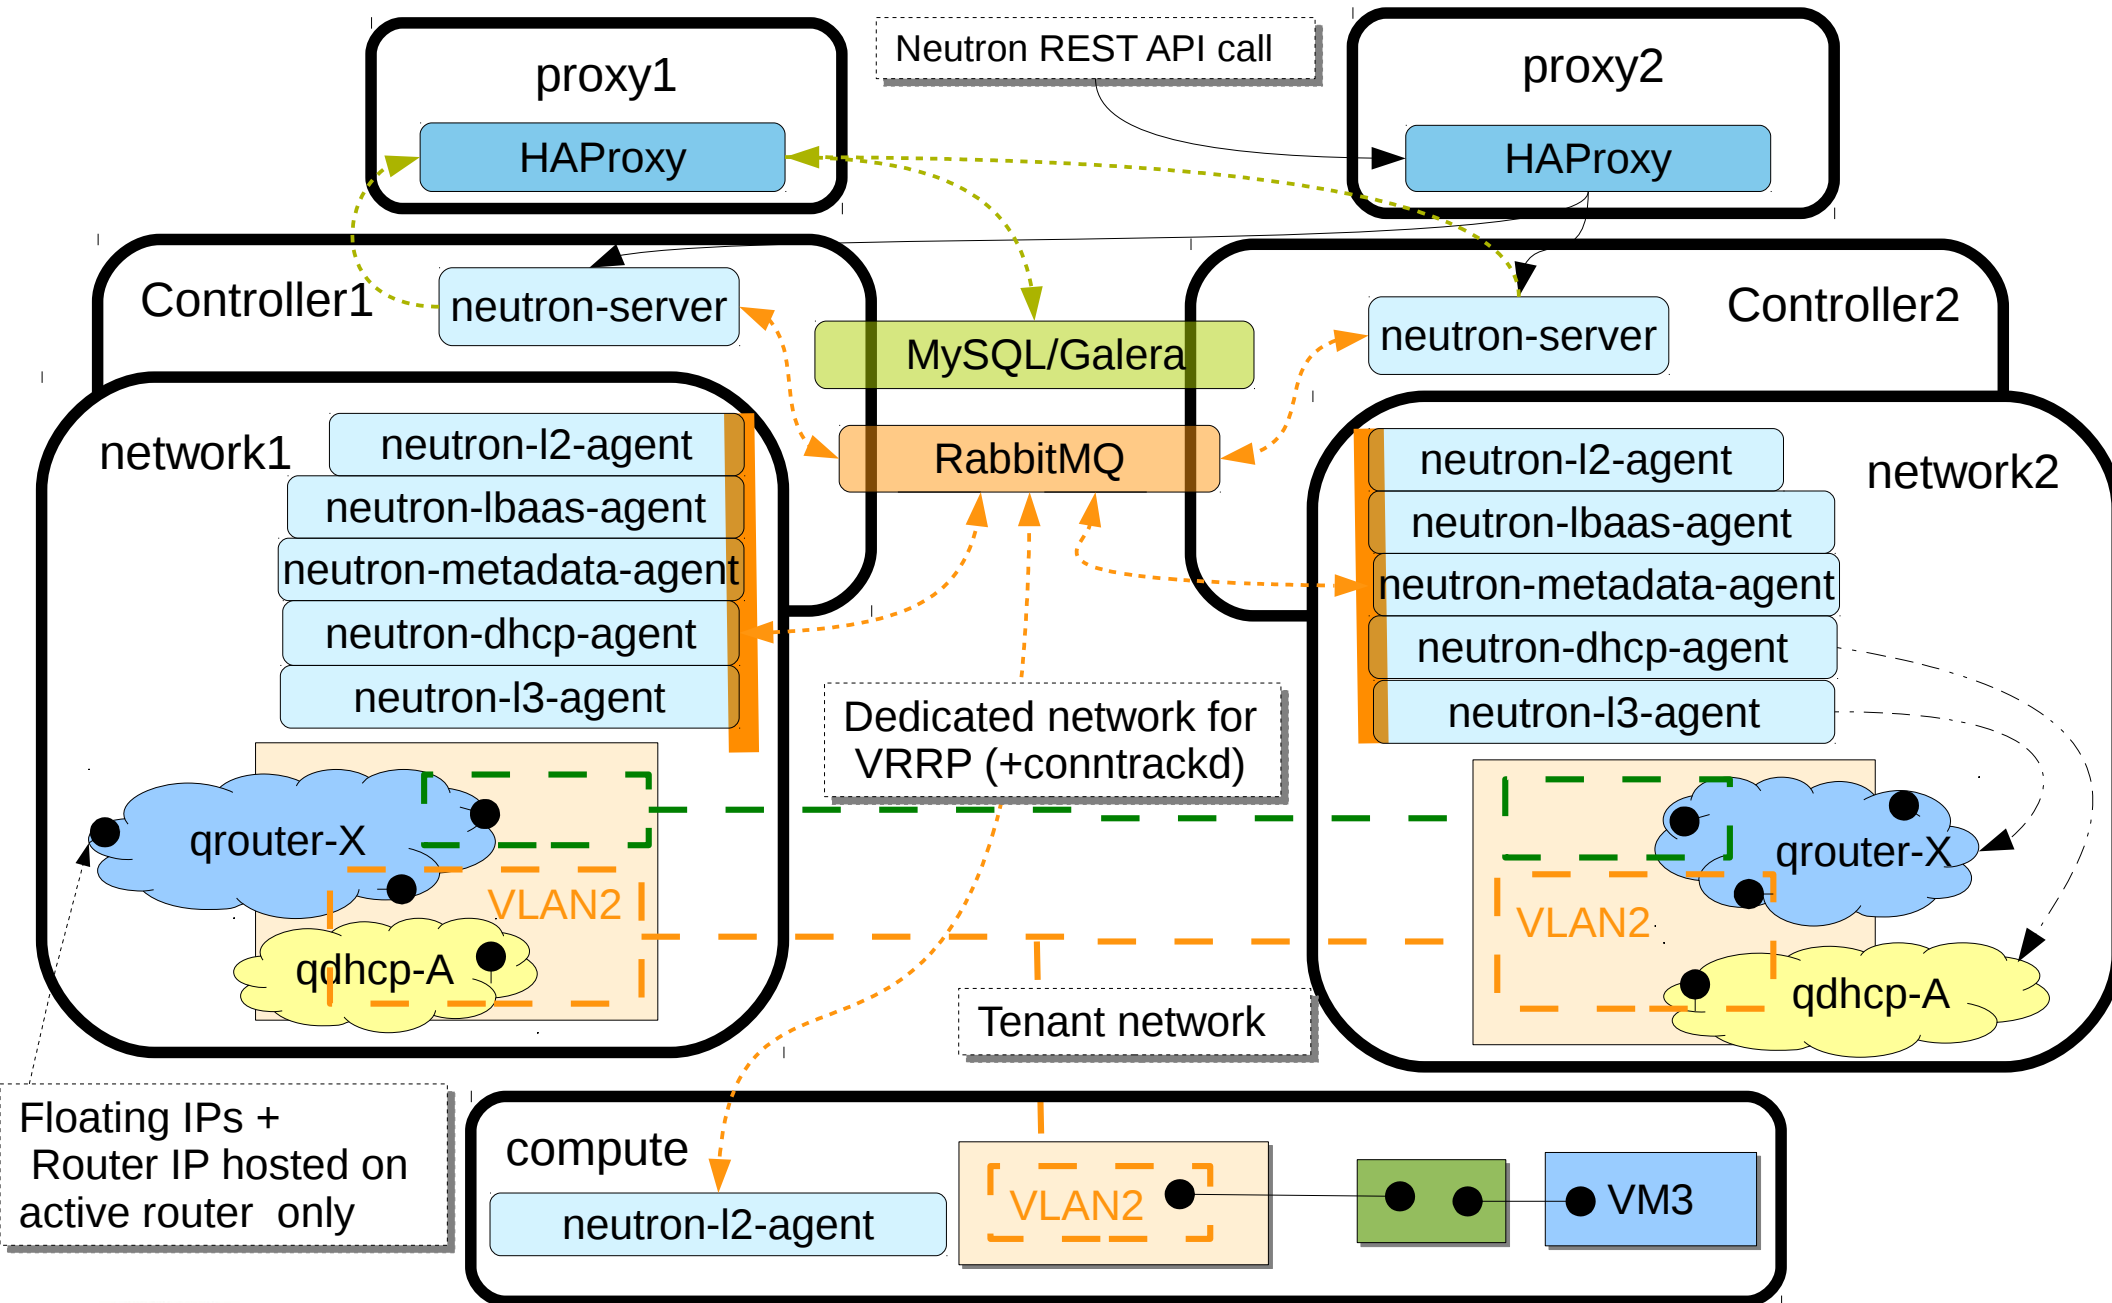





# Ceilometer HA (active-active)



# Neutron - L3 HA



# DVR in Mitaka



# Takeaways

- Nova – no end-to-end solution for VM HA
  - [libvirt-watchdog](#)
  - [proposal for fencing \(nova,neutron,cinder\)](#) (not implemented)
  - vendors provide pacemaker-based failover/fencing for your precious VMs
    - (pacemaker\_remote, compute fence, evacuate)
- Neutron
  - Still “1213 Deadlock found...” bugs with active/active Galera.
  - DVR + HA routers are supported since Mitaka
    - [adv-config-dvr-ha-snat.html](#)
  - LBaaSv2/Octavia implements active/standby LB (haproxy/keepalived)
    - [active\\_passive\\_loadbalancer.html](#)
- Cinder
  - cinder-volume HA has not completed yet
    - [specs/newton/ha-aa-manager\\_locks.html](#)
    - [specs/newton/ha-aa-job-distribution.html](#)



# Questions?

